Zeppelin Notebook Job

Search

ys364_nyu_edu

Funding Safety Analysis

default

```
%spark
var df = spark.read.parquet("/user_data/funding_safety.parquet")

df: org.apache.spark.sql.DataFrame = [School_BEDS_Code: string, Year: int ... 26 more fields]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
df.printSchema()
```

```
root
 |-- School_BEDS_Code: integer (nullable = true)
 |-- Year: integer (nullable = true)
 |-- County_Name: string (nullable = true)
 |-- District_Name: string (nullable = true)
 |-- School_Name: string (nullable = true)
 |-- Homicide: integer (nullable = true)
 |-- Sexual_Offense: integer (nullable = true)
 |-- Assault: integer (nullable = true)
 |-- Weapons_Possession: integer (nullable = true)
 |-- Dignity_Act_Excluding_Cyberbullying: integer (nullable = true)
 |-- Dignity_Act_Cyberbullying: integer (nullable = true)
 |-- Bomb_Threat: integer (nullable = true)
 |-- False_Alarm: integer (nullable = true)
 |-- Drugs: integer (nullable = true)
 |-- Total_Incidents: integer (nullable = true)
 |-- Total_Enrollment: decimal(10,2) (nullable = true)
 |-- Total_Teachers: decimal(10,2) (nullable = true)
 |-- Total_Staff: decimal(10,2) (nullable = true)
 |-- Total_Staff_Enrollment: decimal(28,2) (nullable = true)
 |-- Staff_per_Pupil: decimal(28,2) (nullable = true)
 |-- Total_Funding: decimal(28,2) (nullable = true)
 |-- Total_Funding_per_Pupil: decimal(28,2) (nullable = true)
 |-- Federal_Funding: decimal(28,2) (nullable = true)
 |-- State_&_Local_Funding: decimal(28,2) (nullable = true)
 |-- State_&_Local_Funding_per_Pupil: decimal(28,2) (nullable = true)
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

```
%spark
// // add the total safety incidents column
// val df = df.withColumn(
//   "total_safety_incidents",
//   col("Homicide") + col("Sexual_Offense") + col("Assault") + col("Weapons_Possession") + col("Dignity_Act_Excluding_Cyberbullying") + col("Dignity_Act_Cyberbullying") + col("Bomb_Threat") + col("False_Alarm") + col("Drugs")
// )

// df.org.apache.spark.sql.DataFrame = [School_BEDS_Code: string, Year: int ... 28 more fields]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

Correlation Inspect

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

Global Level

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

```
%spark
val cols = Array("total_safety_incidents", "Total_Enrollment", "Total_Teachers", "Total_Staff", "Total_Funding", "Staff_per_Pupil", "Total_Funding_per_Pupil", "incidents_per_student")
val global_correlations = cols.flatMap { col1 =>
  cols.map { col2 =>
    val corr = df.stat.corr(col1, col2)
    (col1, col2, corr)
  }
}

cols: Array[String] = Array(total_safety_incidents, Total_Enrollment, Total_Teachers, Total_Staff, Total_Funding, Staff_per_Pupil, Total_Funding_per_Pupil, incidents_per_student)
global_correlations: Array[(String, String, Double)] = Array((total_safety_incidents,Total_Enrollment,0.8578385974918845), (total_safety_incidents,Total_Teachers,0.8178978729526), (total_safety_incidents,Total_Staff,0.3661749677295042), (total_safety_incidents,Total_Funding,0.4853578586822456), (total_safety_incidents,Staff_per_Pupil,0.0019779388402242), (total_safety_incidents,Total_Funding_per_Pupil,0.0019779388402242), (total_safety_incidents,incidents_per_student,0.6359941668944885), (Total_Enrollment,total_safety_incidents,0.8578385974918845), (...
```

```
%spark
val corrDF = global_correlations.map {
  case (col1, col2, corr) => (col1, col2, if (corr.isNaN) 0.0 else corr)
}.toDS().toDF("info1", "info2", "correlation")

z.show(corrDF)
```

info1	info2	correlation
Total_Staff	total_safety_incidents	0.3661749677295042
Total_Staff	Total_Enrollment	0.9351667599972648
Total_Staff	Total_Teachers	0.9034566893772301
Total_Staff	Total_Funding	0.8062168936600401
Total_Staff	Staff_per_Pupil	0.32389902080797224
Total_Staff	Total_Funding_per_Pupil	0.27936811489294797
Total_Staff	incidents_per_student	-0.004079519603539963

corrDF: org.apache.spark.sql.DataFrame = [info1: string, info2: string ... 1 more field]

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

Country Level

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

```
%spark
val countyMetrics = df.groupBy("County_Name")
  .agg(
    avg("incidents_per_student").alias("avg_incidents_per_student"),
    avg("Total_Funding_per_Pupil").alias("avg_funding_per_student"),
    avg("Total_Funding").alias("avg_funding"),
    avg("Staff_per_Pupil").alias("avg_staff_ratio"),
    avg("Total_Staff_Enrollment").alias("avg_staff_enrollment"),
    sum("total_safety_incidents").alias("total_safety_incidents"),
    count("School_BEDS_Code").alias("total_county_schools")
  )

countyMetrics: org.apache.spark.sql.DataFrame = [County_Name: string, avg_incidents_per_student: decimal(38,20) ... 6 more fields]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
val columnsToCorrelate = countyMetrics.columns.filter(_ != "County_Name").toSeq
val correlations = columnsToCorrelate.map { col1 =>
  columnsToCorrelate.filter(_ != col1).map { col2 =>
    val corr = countyMetrics.stat.corr(col1, col2).alias("correlation").first().getAs[Double](0)
    (col1, col2, corr)
  }
}.flatMap(identity)

correlations: Array[(String, String, Double)] = Array((avg_incidents_per_student,avg_funding_per_student,0.5265058639049728), (avg_incidents_per_student,avg_teacher_ratio,0.5265058639049728), (avg_incidents_per_student,avg_staff_ratio,0.5148183117294758), (avg_incidents_per_student,avg_funding,0.5148183117294754), (avg_incidents_per_student,avg_staff_enrollment,0.5148183117294754), (avg_incidents_per_student,total_safety_incidents,0.20449731958900803), (avg_incidents_per_student,total_county_schools,0.20449731958900803), (avg_funding_per_student,avg_incidents_per_student,0.5265058639049728), (avg_funding_per_student,avg_teacher_ratio,0.9999999999999999), (avg_funding_per_student,avg_staff_ratio,0.6270436885249674), (avg_incidents_per_student,total_county_enrollment,0.17379123859427115), (avg_incidents_per_student,total_county_schools,0.8944786833313425), (avg_funding_per_student,avg_incidents_per_student,0.5168161117294754), (...
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
// create correlation of from previous map values
val corrDF = correlations.map {
  case (col1, col2, corr) => (col1, col2, if (corr.isNaN) 0.0 else corr)
}.toDS().toDF("info1", "info2", "correlation")

z.show(corrDF)
```

info1	info2	correlation
avg_teacher_ratio	avg_funding_per_student	0.5265058639049728
avg_funding_per_student	avg_teacher_ratio	0.5265058639049728
avg_funding_per_student	avg_incidents_per_student	0.5168183117294758
avg_incidents_per_student	avg_funding_per_student	0.5168183117294754
avg_incidents_per_student	total_county_incidents	0.20449731958900803
total_county_incidents	avg_incidents_per_student	0.20449731958900803
avg_incidents_per_student	total_county_enrollment	0.17379123859427115
total_county_enrollment	avg_incidents_per_student	0.17379123859427115

corrDF: org.apache.spark.sql.DataFrame = [info1: string, info2: string ... 1 more field]

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

Compute Incident Index

Here we use (total incidents count / enrolled students) as the index

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
val schoolTypeAnalysis = df
  .groupBy("School_Type")
  .agg(
    sum("total_safety_incidents").alias("incidents"),
    sum("Total_Funding_per_Pupil").alias("Avg Funding")
  )

schoolTypeAnalysis: org.apache.spark.sql.DataFrame = [School_Type: string, Incidents: double ... 1 more field]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
z.show(schoolTypeAnalysis.orderBy(desc("Incidents")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
z.show(schoolTypeAnalysis.orderBy(desc("Avg Funding")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

Aera Analysis

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

```
// creat analysis
val countyTotalFundingAnalysis = df
  .groupBy("County_Name")
  .agg(sum("Total_Funding"), avg("Total_Funding_per_Pupil"))

countyTotalFundingAnalysis: org.apache.spark.sql.DataFrame = [County_Name: string, sum(Total_Funding): decimal(38,2) ... 1 more field]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
z.show(countyTotalFundingAnalysis.orderBy(desc("sum(Total_Funding)")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
z.show(countyTotalFundingAnalysis.orderBy(desc("Total_Funding_per_Pupil")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
val countySafetyAnalysis = df
  .groupBy("County_Name")
  .agg(
    avg("incidents_per_student").alias("avg_incidents_per_student"),
    sum("total_safety_incidents").alias("total_incidents")
  )

countySafetyAnalysis: org.apache.spark.sql.DataFrame = [County_Name: string, avg_incidents_per_student: decimal(38,20) ... 1 more field]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
z.show(countyTotalFundingAnalysis.orderBy(desc("avg_incidents_per_student")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
z.show(countySafetyAnalysis.orderBy(desc("total_incidents")))
```


Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

Long Term Trend Analysis

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
import org.apache.spark.sql.expressions.Window
val windowSchool = Window.partitionBy("School_BEDS_Code").orderBy("Year")
val schoolChanges = df
  .withColumn("total_incidents_change", lag("incidents_per_student", 1).over(windowSchool).cast("int"))
  .withColumn("total_funding_per_pupil", lag("Total_Funding_per_Pupil", 1).over(windowSchool).cast("double"))
  .withColumn("incident_change", col("incidents_per_student") - col("total_incidents_change"))
  .withColumn("funding_change", col("total_funding_per_pupil") - col("total_funding_per_pupil").cast("double"))
  .withColumn("incidents_change_pct", col("incident_change") / col("total_incidents_change"))
  .withColumn("funding_change_pct", col("funding_change") / col("total_funding_per_pupil").cast("double"))

import org.apache.spark.sql.expressions.Window
windowSchool: org.apache.spark.sql.expressions.WindowSpec = org.apache.spark.sql.expressions.WindowSpec@1738d8c
schoolChanges: org.apache.spark.sql.DataFrame = [School_BEDS_Code: string, Year: int ... 34 more fields]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
val overallChanges = schoolChanges
  .groupBy("School_BEDS_Code", "School_Name", "District_Name", "County_Name")
  .agg(
    max("Year").alias("last_year"),
    min("Year").alias("first_year"),
    sum("incidents_per_student").alias("min_incidents"),
    sum("incidents_per_student").alias("max_incidents"),
    sum("total_funding_per_pupil").alias("min_funding"),
    sum("total_funding_per_pupil").alias("max_funding"),
    sum("incidents_change").alias("total_incidents_change"),
    sum("funding_change").alias("total_funding_change")
  )

overallChanges: org.apache.spark.sql.DataFrame = [School_BEDS_Code: string, School_Name: string ... 12 more fields]
```

Took 0 sec. Last updated by ys364_nyu_edu at December 13 2019, 4:03:42 AM.

finished

```
%spark
val topSchoolChangeSchools = overallChanges.orderBy(desc("total_incident_change")).limit(10).select("School_Name", "District_Name", "County_Name", "total_incident_change", "total_funding_change", "avg_incidents", "avg_funding")
z.show(topSchoolChangeSchools)
```